

FASTER PHP APPS USING

QUEUES & WORKERS

RICHARD BAKER




twitter.com/r_bake_r



github.com/rjbaker



uk.linkedin.com/in/richjbaker




BACKGROUND

MOBILE API

- ▶ 140+ apps
- ▶ ~4.5m users
- ▶ REST/JSON
- ▶ Apache + PHP + ElasticSearch
- ▶ 20+ servers / various tasks
- ▶ Hosted on AWS


THE PROBLEM

TRAFFIC



THE PROBLEM

LATENCY GROWS



**ELIMINATE EXPENSIVE, TIME
CONSUMING OPERATIONS AND
CALLS TO EXTERNAL SERVICES**

USE MYSQL

WHAT TO QUEUE

MYSQL AS A QUEUE

job_id	job_name	data	status
1	sendEmail	{payload}	completed
2	sendPush	{payload}	processing
3	requestThing	{payload}	waiting
4	resizeSelfie	{payload}	waiting

WHAT TO QUEUE

WHY NOT USE TRANSACTIONS TO OBTAIN A LOCK?



**THE DATABASE IS NOT A
QUEUE. THE DATABASE IS
NOT A QUEUE.**

Stephen Corona

USE A QUEUE!

MESSAGE QUEUES VS JOB QUEUES

- ▶ Kind of similar
- ▶ Most job queues built upon some kind of message queue
- ▶ Broker messages between systems
- ▶ Provide transport, storage and protocol
- ▶ Job queues abstract the lower level message component
- ▶ Integrate with most applications fairly easily

QUEUES & WORKERS

JOBS VS SCHEDULED TASKS

- ▶ Run at a predefined point in time 🕒
- ▶ May repeat at regular intervals or according to calendar 📅
- ▶ Typically triggered by cron or other scheduler 🔫
- ▶ Scheduled tasks can trigger jobs! 💥

**CHOOSE A QUEUE WITH
FEATURES BEST SUITED TO
YOUR APPLICATION**

CHOOSING A JOB QUEUE

CONSIDERATIONS

- ▶ Job priority & time sensitivity
- ▶ Job ordering and consistency (FIFO)
- ▶ Payload size limits
- ▶ Message data type & protocol
- ▶ Support for other languages / client libraries
- ▶ Failure management / retry policy
- ▶ Fault tolerance & redundancy
- ▶ One-time delivery guarantee
- ▶ Monitoring & statistics
- ▶ Distribution by task to specific workers e.g. Video encoding

CHOOSING A JOB QUEUE

BEANSTALKD

- ▶ Protocol similar to Memcached
- ▶ Clients need to know about all Beanstalkd servers (like memcached!)
- ▶ Beanstalkd servers can persist jobs, handle restarts without losing jobs
- ▶ Uses “tubes” to differentiate different queues
- ▶ Supports job TTR. Failed/hung jobs get put back into queue.
- ▶ Supports blocking. Client connects and waits for a new job.
- ▶ Requires setup and maintenance
- ▶ Loads of client libraries

<http://kr.github.io/beanstalkd/>

CHOOSING A JOB QUEUE

AMAZON SQS

- ▶ SAAS - Already using AWS, literally no setup
- ▶ Massively redundant, cheap, maintenance free
- ▶ HTTP/JSON under the hood, simple
- ▶ Supports long-polling.
- ▶ Best effort FIFO (no guarantees)
- ▶ No concept of job priority. Use different queues.
- ▶ Retry policy allows jobs to reappear in queue if not completed in specified time
- ▶ Configurable number of retries
- ▶ Queue stats and alarms integrate with autoscaling
- ▶ Scale worker instances based on queue length/backlog/rate

<https://aws.amazon.com/sqs/>

CHOOSING A JOB QUEUE

OTHER POPULAR QUEUES

- ▶ Celery (backed by RabbitMQ) - <http://www.celeryproject.org>
- ▶ php-resque (backed by Redis) - <https://github.com/chrisboulton/php-resque>
- ▶ Kafka - <http://kafka.apache.org>
- ▶ Gearman - <http://gearman.org>
- ▶ Iron.io (SAAS) - <https://www.iron.io>
- ▶ Loads more - www.queues.io

PROCESSING JOBS

PROCESSING JOBS

WORKER PROCESS

- ▶ Essentially an infinite loop
- ▶ Executed on command line
- ▶ Asks queue for new job
- ▶ Resolves job method
- ▶ Execute with payload
- ▶ Delete job from queue
- ▶ Repeat

PROCESSING JOBS

```
<?php
```

```
$queue = new Queue();
```

```
while(true) {
```

```
    $job = $queue->pop('queue-name');
```

```
    try {
```

```
        if ($job->execute()) {
```

```
            $job->delete();
```

```
        } else {
```

```
            $job->release();
```

```
        }
```

```
    } catch (\Exception $e) {
```

```
        $job->release();
```

```
    }
```

```
}
```

IMPROVING THE WORKER


```
pcntl_signal_dispatch();
```

PROCESS CONTROL EXTENSIONS (PCNTL)

- ▶ Respond to unix process signals
- ▶ Gracefully stop worker processes
- ▶ Complete current job before exiting
- ▶ Careful if using Apache mod_php on same server
- ▶ <http://php.net/manual/en/book.pcntl.php>

IMPROVED WORKER

```
<?php

namespace Demo;

use Demo\Queue\QueueInterface;

class Worker
{
    protected $shouldRun = true;

    protected $queue;

    public function __construct(QueueInterface $queue)
    {
        declare(ticks = 1);

        $this->queue = $queue;

        pcntl_signal(SIGTERM, [$this, 'signalHandler']);
        pcntl_signal(SIGINT, [$this, 'signalHandler']);
        pcntl_signal(SIGQUIT, [$this, 'signalHandler']);
    }
}
```

```

public function run($queueName)
{
    echo "Starting worker on queue '{$queueName}' \n";

    while ($this->shouldRun) {

        $job = $this->queue->pop($queueName);

        try {
            if ($job->execute()) {
                $job->delete();
            } else {
                $job->release();
            }
        } catch (\Exception $e) {
            $job->release();
            error_log($e->getTraceAsString());
        }

        pcntl_signal_dispatch();
    }
}

public function signalHandler($signal)
{
    switch ($signal) {
        case SIGTERM:
        case SIGINT:
        case SIGQUIT:
            echo "Job completed. Exiting... \n";
            $this->shouldRun = false;
            break;
    }
}
}

```

WTF IS `DECLARE(TICKS = 1);?`

- ▶ Officially deprecated
- ▶ Triggered after php has executed a certain number of statements
- ▶ Interacts with `pcntl_signal_dispatch()`
- ▶ I admit i've not fully tested this with PHP7.0

PROCESSING JOBS

KEEPING WORKERS RUNNING



PROCESSING JOBS

SUPERVISOR

- ▶ Process manager in similar vein to forever, pm2, php-fpm
- ▶ Runs as service
- ▶ Starts and restarts php worker processes
- ▶ Has CLI client (supervisorctl)
- ▶ Web interface
- ▶ Easy to install and configure

<http://supervisord.org>

PROCESSING JOBS

SUPERVISOR CONFIG

```
[program:alertworker]
command = /usr/bin/php /path/to/queueRunner.php -q=prod-alerts
autorestart = true
autostart = true
directory = /path/to/scripts
environment = DEPLOYMENT='production'
exitcodes = 0,2
numprocs = 1
numprocs_start = 0
priority = 999
startretries = 3
startsecs = 4
stderr_capture_maxbytes = 1MB
stderr_events_enabled = false
stderr_logfile = AUTO
stderr_logfile_backups = 10
stderr_logfile_maxbytes = 50MB
stderr_syslog = false
stdout_capture_maxbytes = 1MB
stdout_events_enabled = true
stdout_logfile = AUTO
stdout_logfile_backups = 10
stdout_logfile_maxbytes = 40MB
stdout_syslog = false
stopsignal = TERM
stopwaitsecs = 10
umask = 022
user = worker
```

DEMO TIME

THANKS FOR LISTENING!



twitter.com/r_bake_r



github.com/rjbaker



uk.linkedin.com/in/richjbaker